

libsemx Manuscript Draft

Deniz Akdemir and contributors

2025-12-29

Contents

0.1	Abstract	1
0.2	Introduction	1
0.3	Example 1: Random slope LMM (sleepstudy, R)	2
0.4	Example 1b: Same model via Python	3
0.5	Example 2: CFA / SEM (BFI personality items, R)	4
0.6	Example 3: Survival regression (ovarian, R)	7
0.7	Example 4: Genomic mixed model (Maize Diversity Panel, R)	8
0.8	Comparisons to lme4 and sommer (runtime + variance components)	9
0.9	Comparisons to related libraries	11
0.10	Reproducibility and next steps	13
0.11	Why libsemx (and why it can be slower on tiny models)	13
0.12	Interpreting fit indices and degrees of freedom	13

0.1 Abstract

libsemx is a unified C++20 engine for structural equation modeling (SEM) and generalized linear mixed models (GLMM) with Python and R bindings. It couples a shared intermediate representation (ModelIR), flexible covariance structures (including genomic kernels), and support for non-Gaussian families and survival outcomes. We illustrate the modeling API, compare runtime/estimates against lme4 (Bates et al. 2015), sommer (Covarrubias-Pazaran 2016), and statsmodels (Seabold and Perktold 2010), and discuss when to prefer libsemx (SEM+mixed, non-Gaussian, custom covariance) versus specialized libraries. The project is openly developed on GitHub and intended for preprint submission (e.g., arXiv).

0.2 Introduction

libsemx is a unified C++20 engine for structural equation modeling (SEM) and generalized linear mixed models (GLMM), exposed through Python and R. It keeps a shared intermediate representation (ModelIR) for variables, paths, covariances, and random effects so both bindings produce identical likelihoods, gradients, and diagnostics.

Architecture at a glance

- Variables: observed, latent, grouping, exogenous; observed nodes carry families (gaussian, binomial, poisson, negative binomial, gamma, Weibull, exponential, log-normal, log-logistic, ordinal).
- Edges: loadings ($=\sim$), regressions (\sim), covariances ($\sim\sim$) tied to parameter IDs for constraints/equality.
- Covariances: unstructured, diagonal, compound symmetry, AR(1), Toeplitz, Kronecker, fixed/scaled fixed, multi-kernel simplex blends.
- Random effects: named effects referencing covariances; optional `lambda` shrinkage for ridge and genomic prediction.
- Estimation: ML/REML with L-BFGS or gradient descent, Average Information REML for variance components, Laplace for non-Gaussian random effects, spectral shortcuts for dense genomic kernels, FIML for missing data, SEM fit indices (2 , CFI, TLI, RMSEA, SRMR).

0.3 Example 1: Random slope LMM (sleepstudy, R)

Real data (`data/sleepstudy.csv`) with subject-specific intercepts and slopes:

```
library(dplyr)
sleep_df <- readr::read_csv("data/sleepstudy.csv", show_col_types = FALSE) %>%
  select(-rownames) %>%
  mutate(Reaction = Reaction / 100)

model_lmm <- semx_model(
  equations = c("Reaction ~ Days + (Days | Subject)"),
  families = c(Reaction = "gaussian")
)

t_semx <- system.time({
  fit_lmm <- semx_fit(model_lmm, sleep_df, optimizer_name = "lbfgs")
})
fit_lmm_elapsed <- t_semx[["elapsed"]]

summary(fit_lmm)           # fixed effects, variances, fit indices

#> Optimization converged: TRUE
#> Iterations: 129
#> Log-likelihood: -47.039
#> Chi-square: 106.078 (df=NA)
#> AIC: 106.1, BIC: 125.2
#>
#>                               Estimate   Std.Error    z.value   P.value
#> beta_Reaction_on_Days      0.104672860 0.015022367 6.9678006 3.219425e-12
#> alpha_Reaction_on_intercept 2.514051048 0.066322768 37.9063046 0.0000000e+00
#> psi_Reaction_Reaction     0.065494103 0.007718554 8.4852813 0.0000000e+00
#> cov_re_1_0                  0.237805670 0.055773710 4.2637592 2.010161e-05
#> cov_re_1_1                  0.004648935 0.018279776 0.2543212 7.992474e-01
#> cov_re_1_2                  0.056979003 0.012291142 4.6357778 3.555979e-06
#>
#> Variance Components:
#>   Group      Name1 Name2   Variance   Std.Dev       Corr
#>   Subject _intercept        0.056551536 0.23780567          NA
#>   Subject _intercept  Days 0.001105543          NA 0.0813201
#>   Subject      Days        0.003268219 0.05716834          NA
```

```

semx_ranef(fit_lmm) %>% head()    # BLUPs per random-effect block

#> $re_Subject_1
#>      _intercept          Days
#> 308  0.028156841  0.090755338
#> 309 -0.400484878 -0.086440674
#> 310 -0.384331534 -0.055133789
#> 330  0.228322947 -0.046587502
#> 331  0.215499891 -0.029445199
#> 332  0.088155864 -0.002352091
#> 333  0.164419872 -0.001588237
#> 334 -0.069967202  0.010327362
#> 335 -0.010374193 -0.105994435
#> 337  0.346663149  0.086323772
#> 349 -0.245581581  0.010644008
#> 350 -0.123346311  0.064717035
#> 351  0.042740664 -0.029553436
#> 352  0.206222197  0.035617042
#> 369  0.032585362  0.008717103
#> 370 -0.247103332  0.046597354
#> 371  0.007232813 -0.009710559
#> 372  0.121189431  0.013106909

```

0.4 Example 1b: Same model via Python

The same IR is built and fit from Python using `reticulate`:

```

import pandas as pd
import numpy as np
from semx import Model

sleep_df = pd.read_csv("data/sleepstudy.csv")
sleep_df["Reaction"] = sleep_df["Reaction"] / 100.0
# Encode grouping as integer codes (required by C++ backend)
sleep_df["Subject"] = pd.Categorical(sleep_df["Subject"]).codes.astype(int)

# Explicit residual variance term and random slope
spec = Model(
    equations=[
        "Reaction ~ 1 + Days + (Days | Subject)",
        "Reaction ~~ Reaction",
    ],
    families={"Reaction": "gaussian", "Days": "gaussian"},
    kinds={"Subject": "grouping"},
)
fit = spec.fit(sleep_df)

param_map = dict(zip(fit.fit_result.parameter_names, fit.fit_result.optimization_result.parameters))
print("Parameters (name=value):")

#> Parameters (name=value):

```

```

for k, v in param_map.items():
    print(f" {k}: {v:.6f}")

#>   beta_Reaction_on__intercept: 2.514051
#>   beta_Reaction_on_Days: 0.104673
#>   psi_Reaction_Reaction: 0.065494
#>   cov_re_1_0: 0.237806
#>   cov_re_1_1: 0.004649
#>   cov_re_1_2: 0.056979

# Manual BLUPs (pybind currently returns zeros for random effects; compute directly)
beta = np.array([param_map["beta_Reaction_on__intercept"], param_map["beta_Reaction_on_Days"]])
sigma_e = param_map["psi_Reaction_Reaction"]
G = np.array([
    [param_map["cov_re_1_0"], param_map["cov_re_1_1"]],
    [param_map["cov_re_1_1"], param_map["cov_re_1_2"]],
])
blup_rows = []
for subj, df_g in sleep_df.groupby("Subject"):
    y = df_g["Reaction"].to_numpy()
    days = df_g["Days"].to_numpy()
    X = np.column_stack([np.ones_like(days), days]) # fixed + random design
    resid = y - X.dot(beta)
    V = X.dot(G).dot(X.T) + sigma_e * np.eye(len(y))
    u = G.dot(X.T).dot(np.linalg.solve(V, resid))
    blup_rows.append({"Subject": subj, "u_intercept": u[0], "u_days": u[1]})

blups = pd.DataFrame(blup_rows).sort_values("Subject")
print("\nBLUPs (first 6 subjects):")

#>
#> BLUPs (first 6 subjects):

print(blups.head(6))

#>   Subject  u_intercept    u_days
#> 0        0     -0.058670  0.110399
#> 1        1     -0.428320 -0.087291
#> 2        2     -0.440213 -0.049516
#> 3        3      0.344710 -0.068275
#> 4        4      0.310120 -0.046555
#> 5        5      0.116699 -0.007110

```

0.5 Example 2: CFA / SEM (BFI personality items, R)

Confirmatory factor analysis on the Big Five inventory (`data/bfi.csv`). Latent variances are fixed to 1 for scale identification.

```

bfi_df <- readr::read_csv("data/bfi.csv", show_col_types = FALSE) %>%
  select(-rownames) %>%
  na.omit() %>%
  slice_head(n = 600)

item_cols <- c(paste0("A", 1:5), paste0("C", 1:5), paste0("E", 1:5), paste0("N", 1:5), paste0("O", 1:5))
bfi_df[item_cols] <- scale(bfi_df[item_cols])

equations <- c(
  "Agreeableness =~ NA*A1 + A2 + A3 + A4 + A5",
  "Conscientiousness =~ NA*C1 + C2 + C3 + C4 + C5",
  "Extraversion =~ NA*E1 + E2 + E3 + E4 + E5",
  "Neuroticism =~ NA*N1 + N2 + N3 + N4 + N5",
  "Openness =~ NA*O1 + O2 + O3 + O4 + O5",
  "Agreeableness ~~ 1*Agreeableness",
  "Conscientiousness ~~ 1*Conscientiousness",
  "Extraversion ~~ 1*Extraversion",
  "Neuroticism ~~ 1*Neuroticism",
  "Openness ~~ 1*Openness"
)

families <- setNames(rep("gaussian", length(item_cols)), item_cols)

# Mildly informative starting values to stabilize L-BFGS
init_params <- list()
for (eq in equations) {
  if (grepl("=~", eq)) {
    parts <- strsplit(eq, "=~")[[1]]
    latent <- trimws(parts[1])
    inds <- trimws(strsplit(parts[2], "\\+"))[[1]]
    for (ind_raw in inds) {
      ind <- gsub("NA\\*", "", trimws(ind_raw))
      init_params[[paste0("lambda_", ind, "_on_", latent)]] <- 0.5
      init_params[[paste0("psi_", ind, "_", ind)]] <- 0.5
    }
  }
}

model_cfa <- semx_model(
  equations = equations,
  families = families,
  parameters = init_params
)

fit_cfa <- semx_fit(model_cfa, bfi_df, optimizer_name = "lbfgs")
summary(fit_cfa)           # loadings, residual variances, fit indices

```

```

#> Optimization converged: TRUE
#> Iterations: 34
#> Log-likelihood: -19816.803
#> Chi-square: 39733.606 (df=300)
#> P-value: 0.000
#> CFI: 0.671, TLI: 0.671, RMSEA: 0.089, SRMR: 0.142

```

	Estimate	Std.Error	z.value	P.value
#> lambda_A1_on_Agreeableness	-0.3363039	0.04668813	-7.203200	5.881962e-13
#> lambda_A2_on_Agreeableness	0.6610674	0.04379622	15.094167	0.000000e+00
#> lambda_A3_on_Agreeableness	0.7153662	0.04306594	16.610952	0.000000e+00
#> lambda_A4_on_Agreeableness	0.4863282	0.04485702	10.841742	0.000000e+00
#> lambda_A5_on_Agreeableness	0.6278129	0.04361900	14.393107	0.000000e+00
#> lambda_C1_on_Conscientiousness	0.4821873	0.04616018	10.445958	0.000000e+00
#> lambda_C2_on_Conscientiousness	0.6434928	0.04560025	14.111606	0.000000e+00
#> lambda_C3_on_Conscientiousness	0.5688728	0.04560263	12.474562	0.000000e+00
#> lambda_C4_on_Conscientiousness	-0.6281092	0.04603912	-13.642946	0.000000e+00
#> lambda_C5_on_Conscientiousness	-0.5268256	0.04667944	-11.286031	0.000000e+00
#> lambda_E1_on_Extraversion	-0.5969958	0.04274728	-13.965701	0.000000e+00
#> lambda_E2_on_Extraversion	-0.7148705	0.04177410	-17.112767	0.000000e+00
#> lambda_E3_on_Extraversion	0.6008709	0.04329340	13.879042	0.000000e+00
#> lambda_E4_on_Extraversion	0.6412587	0.04212011	15.224525	0.000000e+00
#> lambda_E5_on_Extraversion	0.5474789	0.04359838	12.557322	0.000000e+00
#> lambda_N1_on_Neuroticism	0.8109395	0.03619465	22.404956	0.000000e+00
#> lambda_N2_on_Neuroticism	0.8346101	0.03587671	23.263286	0.000000e+00
#> lambda_N3_on_Neuroticism	0.7297904	0.03820746	19.100730	0.000000e+00
#> lambda_N4_on_Neuroticism	0.5120498	0.04198083	12.197227	0.000000e+00
#> lambda_N5_on_Neuroticism	0.5223308	0.04120977	12.674927	0.000000e+00
#> lambda_O1_on_Openness	0.5377736	0.04733319	11.361448	0.000000e+00
#> lambda_O2_on_Openness	-0.4039174	0.04914854	-8.218298	2.220446e-16
#> lambda_O3_on_Openness	0.7047112	0.05009588	14.067250	0.000000e+00
#> lambda_O4_on_Openness	0.3442046	0.04870034	7.067807	1.574074e-12
#> lambda_O5_on_Openness	-0.5396297	0.04986444	-10.821935	0.000000e+00
#> psi_A1_A1	0.8852330	0.05374155	16.472040	0.000000e+00
#> psi_A2_A2	0.5613232	0.04616556	12.158917	0.000000e+00
#> psi_A3_A3	0.4865845	0.04530684	10.739758	0.000000e+00
#> psi_A4_A4	0.7618182	0.04950252	15.389482	0.000000e+00
#> psi_A5_A5	0.6041843	0.04594582	13.149931	0.000000e+00
#> psi_C1_C1	0.7658288	0.05068899	15.108383	0.000000e+00
#> psi_C2_C2	0.5842504	0.04896502	11.931996	0.000000e+00
#> psi_C3_C3	0.6747170	0.04903766	13.759160	0.000000e+00
#> psi_C4_C4	0.6038121	0.04953052	12.190709	0.000000e+00
#> psi_C5_C5	0.7207881	0.05060029	14.244743	0.000000e+00
#> psi_E1_E1	0.6419294	0.04505546	14.247538	0.000000e+00
#> psi_E2_E2	0.4872935	0.04275529	11.397268	0.000000e+00
#> psi_E3_E3	0.6372875	0.04574054	13.932665	0.000000e+00
#> psi_E4_E4	0.5871206	0.04351983	13.490877	0.000000e+00
#> psi_E5_E5	0.6986002	0.04701422	14.859337	0.000000e+00
#> psi_N1_N1	0.3407105	0.02996233	11.371297	0.000000e+00
#> psi_N2_N2	0.3017593	0.02951567	10.223697	0.000000e+00
#> psi_N3_N3	0.4657393	0.03512550	13.259293	0.000000e+00
#> psi_N4_N4	0.7361384	0.04624708	15.917510	0.000000e+00
#> psi_N5_N5	0.7255039	0.04516783	16.062403	0.000000e+00
#> psi_O1_O1	0.7091329	0.05120514	13.848861	0.000000e+00
#> psi_O2_O2	0.8351841	0.05389234	15.497268	0.000000e+00
#> psi_O3_O3	0.5017154	0.05780255	8.679815	0.000000e+00
#> psi_O4_O4	0.8798565	0.05443095	16.164635	0.000000e+00
#> psi_O5_O5	0.7071331	0.05392201	13.113997	0.000000e+00

```
# Path diagram (optional; not exported in NAMESPACE, use ::::)
# semx:::semx_plot_path(fit_cfa)
```

0.6 Example 3: Survival regression (ovarian, R)

Weibull survival with censoring (`data/ovarian_survival.csv`), fixed effects on age and treatment arm:

```
ovarian <- readr::read_csv("data/ovarian_survival.csv", show_col_types = FALSE) %>%
  select(-rownames)

surv_model <- semx_model(
  equations = c("Surv(futime, fustat) ~ age + rx + ecog.ps"),
  families = c(futime = "weibull", age = "gaussian", rx = "gaussian", ecog.ps = "gaussian")
)

fit_surv <- semx_fit(
  surv_model,
  ovarian,
  options = list(max_iterations = 300, tolerance = 1e-5, force_laplace = TRUE),
  optimizer_name = "lbfgs"
)

summary(fit_surv) # regression coefficients, scale/shape
```

#> Optimization converged: TRUE
#> Iterations: 6
#> Log-likelihood: -327.766
#> Chi-square: 671.532 (df=6)
#> P-value: 0.000
#> CFI: -0.885, TLI: -0.885, RMSEA: 0.394, SRMR: 10.911
#> AIC: 671.5, BIC: 681.6
#>

	Estimate	Std.Error	z.value	P.value
#> beta_futime_on_age	-0.07965424	0.01999171	-3.9843643	6.766100e-05
#> beta_futime_on_rx	0.56114462	0.34087855	1.6461717	9.972842e-02
#> beta_futime_on_ecog.ps	0.06019814	0.33114847	0.1817859	8.557507e-01
#> alpha_futime_on_intercept	10.40851489	1.46333168	7.1128884	1.136424e-12
#> psi_futime_futime	1.82751637	0.43298116	4.2207757	2.434631e-05
#> psi_age_age	3252.65051220	902.12294613	3.6055512	3.114910e-04
#> psi_rx_rx	2.49999983	0.69337517	3.6055514	3.114908e-04
#> psi_ecog.ps_ecog.ps	2.38461539	0.66137331	3.6055513	3.114910e-04

Predict survival at selected times

```
semx_predict_survival(fit_surv, newdata = ovarian[1:3, ], times = c(100, 300, 500), outcome = "futime")
```

#> 100 300 500
#> 1 0.7425899 0.10902957 0.0035644391
#> 2 0.6651962 0.04804094 0.0004432768
#> 3 0.8926935 0.42944693 0.1164941986

0.7 Example 4: Genomic mixed model (Maize Diversity Panel, R)

Single-kernel GBLUP on standardized height (EarHT) using SNP markers (data/mdp_numeric.csv, data/mdp_traits.csv). To keep knitting fast, a small marker subset is used; increase p_subset for fuller analyses.

```

library(Matrix)

numeric_df <- readr::read_csv("data/mdp_numeric.csv", show_col_types = FALSE)
traits_df <- readr::read_csv("data/mdp_traits.csv", show_col_types = FALSE) %>%
  rename(taxa = Taxa)

merged_df <- numeric_df %>%
  inner_join(traits_df, by = c("taxa" = "taxa")) %>%
  filter(!is.na(EarHT))

marker_cols <- setdiff(names(numeric_df), "taxa")
p_subset <- 300
marker_cols <- marker_cols[seq_len(min(p_subset, length(marker_cols)))] 

M <- as.matrix(merged_df[, marker_cols])
merged_df$EarHT_std <- scale(merged_df$EarHT)
merged_df$all_groups <- 1 # single grouping indicator for random effect

model_gblup <- semx_model(
  equations = c("EarHT_std ~ 1"),
  families = c(EarHT_std = "gaussian"),
  genomic = list(polygenic = list(markers = M, structure = "grm")),
  random_effects = list(
    list(name = "u", variables = c("all_groups"), covariance = "polygenic")
  )
)

fit_gblup <- semx_fit(model_gblup, merged_df, optimizer_name = "lbfgs")
summary(fit_gblup) # genetic vs residual variance

#> Optimization converged: TRUE
#> Iterations: 8
#> Log-likelihood: -365.835
#> Chi-square: 737.670 (df=NA)
#> AIC: 737.7, BIC: 748.6
#>
#>                               Estimate Std.Error      z.value      P.value
#> alpha_EarHT_std_on__intercept 3.242287e-10 0.04460759 7.268465e-09 1.000000e+00
#> psi_EarHT_std_EarHT_std       5.551634e-01 0.06747569 8.227607e+00 2.220446e-16
#> polygenic_0                   3.819819e-01 0.10538957 3.624475e+00 2.895486e-04
#>
#> Variance Components:
#>      Group      Name1 Name2 Variance Std.Dev Corr
#> all_groups (Intercept)   0.3819829 0.6180476   NA

# Extract variance components and heritability (h2)
param_names <- fit_gblup$parameter_names

```

```

if (is.null(param_names) || length(param_names) == 0) {
  # Fallback to IR parameter ids
  param_names <- fit_gblup$model$ir$parameter_ids()
}
params <- setNames(fit_gblup$optimization_result$parameters, param_names)

genetic_key <- "polygenic_0"
if (is.null(params[[genetic_key]])) {
  poly_keys <- grep("polygenic", names(params), value = TRUE)
  if (length(poly_keys)) genetic_key <- poly_keys[[1]]
}
resid_key <- "psi_EarHT_std_EarHT_std"

genetic_var <- params[[genetic_key]]
resid_var <- params[[resid_key]]

if (!is.null(genetic_var) && !is.null(resid_var)) {
  h2 <- genetic_var / (genetic_var + resid_var)
  vc_tbl <- data.frame(
    component = c("Genetic (GRM)", "Residual"),
    variance = c(genetic_var, resid_var),
    proportion = c(genetic_var, resid_var) / (genetic_var + resid_var)
  )
  print(vc_tbl)
  cat(sprintf("Narrow-sense heritability (h2): %.3f\n", h2))
} else {
  cat(
    "Could not locate variance components to compute h2; available parameters:\n",
    paste(names(params), collapse = ", "),
    "\n"
  )
}
}

#>      component variance proportion
#> 1  Genetic (GRM) 0.3819819  0.4076015
#> 2      Residual 0.5551634  0.5923985
#> Narrow-sense heritability (h2): 0.408

```

0.8 Comparisons to lme4 and sommer (runtime + variance components)

Small side-by-side fits on shared datasets; guarded by `requireNamespace()` to keep knitting robust.

```

library(tibble)

compare_rows <- list()

# libsemx (sleepstudy random slopes; already fit above)
if (exists("fit_lmm")) {
  pe <- fit_lmm$optimization_result$parameters
  pn <- fit_lmm$parameter_names
  if (!is.null(pn) && length(pn) == length(pe)) {
    pe <- setNames(pe, pn)
  }
}

```

```

}

beta0 <- pe[["alpha_Reaction_on__intercept"]] %||% pe[["beta_Reaction_on__intercept"]] %||% pe[[1]]
beta1 <- pe[["beta_Reaction_on_Days"]] %||% pe[[2]]
# Random-effect Cholesky diagonals -> variances
var_intercept <- pe[["cov_re_1_0"]] %||% pe[[length(pe) - 2]]
var_slope <- pe[["cov_re_1_2"]] %||% pe[[length(pe) - 0]]
if (!is.null(names(pe))) {
  if ("cov_re_1_0" %in% names(pe)) var_intercept <- var_intercept^2
  if ("cov_re_1_2" %in% names(pe)) var_slope <- var_slope^2
}
compare_rows[["libsemx_sleepstudy"]] <- tibble(
  package = "libsemx",
  model = "sleepstudy",
  beta0 = beta0,
  beta1 = beta1,
  var_intercept = var_intercept,
  var_slope = var_slope,
  elapsed = fit_lmm_elapsed %||% NA_real_
)
}

# lme4 on sleepstudy
if (requireNamespace("lme4", quietly = TRUE)) {
  t_lme4 <- system.time({
    lme4_fit <- lme4::lmer(Reaction ~ Days + (Days | Subject), data = sleep_df)
  })
  vc_mat <- as.matrix(lme4::VarCorr(lme4_fit)$Subject)
  intercept_var <- vc_mat[1, 1]
  slope_var <- vc_mat[2, 2]
  compare_rows[["lme4_sleepstudy"]] <- tibble(
    package = "lme4",
    model = "sleepstudy",
    beta0 = lme4::fixef(lme4_fit)[["(Intercept)"]],
    beta1 = lme4::fixef(lme4_fit)[["Days"]],
    var_intercept = intercept_var,
    var_slope = slope_var,
    elapsed = t_lme4[["elapsed"]]
  )
}

# sommer on MDP GBLUP (same marker subset as above)
if (requireNamespace("sommer", quietly = TRUE)) {
  Gmat <- tcrossprod(scale(M)) / ncol(M)
  rownames(Gmat) <- merged_df$taxa
  colnames(Gmat) <- merged_df$taxa
  mdp_dat <- merged_df %>% mutate(taxa = factor(taxa))
  sommer_res <- tryCatch({
    t_sommer <- system.time({
      sommer_fit <- sommer::mmer(
        EarHT_std ~ 1,
        random = ~ sommer::vsr(taxa, Gu = Gmat),
        data = mdp_dat
    })
  })
}

```

```

    })
  list(fit = sommer_fit, elapsed = t_sommer[["elapsed"]])
}, error = function(e) {
  message("sommer failed: ", conditionMessage(e))
  NULL
})
if (!is.null(sommer_res)) {
  vc_sommer <- sommer::summary.mmer(sommer_res$fit)$varcomp
  compare_rows[["sommer_mdp"]] <- tibble(
    package = "sommer",
    model = "mdp_gblup",
    beta0 = sommer_res$fit$Beta[1, 1],
    beta1 = NA_real_,
    var_intercept = vc_sommer["u:taxa", "VarComp"],
    var_slope = NA_real_,
    elapsed = sommer_res$elapsed
  )
}
}

if (length(compare_rows)) {
  dplyr::bind_rows(compare_rows)
} else {
  message("No comparison packages available.")
}
}

#> # A tibble: 2 x 7
#>   package model      beta0 beta1 var_intercept var_slope elapsed
#>   <chr>   <chr>     <dbl> <dbl>      <dbl>     <dbl>    <dbl>
#> 1 libsemx sleepstudy  2.51  0.105      0.0566    0.00325   1.36
#> 2 lme4    sleepstudy  2.51  0.105      0.0612    0.00351   0.463

```

0.9 Comparisons to related libraries

Library	Scope	Families	Covariance/kernels	SEM support	Notes
libsemx	SEM + GLMM unified	Gaussian, GLM, survival, ordinal	Unstructured, CS, AR(1), Toeplitz, Kronecker, genomic/multi-kernel	Yes	C++ core, Python/R front-ends, spectral short-cuts, FIML

Library	Scope	Families	Covariance/kernels	SEM support	Notes
statsmodels MixedLM	Linear mixed models	Mostly Gaussian	Random intercept/slope, limited structures	No	No Laplace for non-Gaussian; no latent variables
lme4	GLMM (R)	Gaussian, binomial, Poisson (others via glmer)	Random effects with simple variance components	Limited	No SEM paths or fit indices; no genomic kernels
nlme	Linear mixed models (R)	Gaussian	Correlation/covariance structures, no kernels	No	Older API, limited non-Gaussian support
lavaan	SEM (R)	Mostly Gaussian	Latent covariance structures	Yes	No mixed-model random effects or GLMM families
sommer	Mixed models with kernels (R)	Gaussian	Genomic kernels, multi-trait	No	Genomic focus; limited non-Gaussian families

0.10 Reproducibility and next steps

- Data live in `data/` and examples mirror `docs/examples/libsemx_starter.Rmd` and `docs/examples/mdp_analysis.Rmd`.
- Python examples: `python/examples/shrinkage_example.py`, `python/examples/crossed_effects_example.py`.
- R example: `Rpkg/semx/examples/shrinkage_example.R`.

Planned manuscript additions: benchmark tables against `lme4/nlme/lavaan/sommer/statsmodels`, multi-group invariance SEM, competing risks survival, and multi-kernel simplex genomic prediction with real kernels.

0.11 Why libsemx (and why it can be slower on tiny models)

- **General engine vs. special-case speed:** `lme4/nlme` have hand-optimized Gaussian REML/ML paths; they are very fast on small linear mixed models. `libsemx` runs a unified L-BFGS/Laplace/FIML stack that can handle non-Gaussian families, latent variables, survival, and custom covariances. That generality adds startup/solver overhead on toy problems (`sleepstudy`), but becomes advantageous as models get richer (ordinal, survival, multi-kernel, SEM).
- **Unified IR and cross-language parity:** Python and R bindings share the same ModelIR and likelihood driver, so estimates/gradients/fit indices match across languages and with the C++ core. This reduces drift between front-ends.
- **Non-Gaussian and survival support:** Built-in Laplace approximation, competing risks/survival families, and dispersion handling go beyond what `lme4/nlme` provide.
- **Flexible covariances and genomics:** Supports AR(1)/Toeplitz/Kronecker, multi-kernel simplex blends, and spectral shortcuts for dense kernels; integrates genomic prediction workflows out of the box.
- **SEM + mixed models together:** `lavaan`-style SEM paths with mixed-model random effects and GLMM families, plus SEM fit indices and modification indices.
- **Diagnostics and post-estimation:** Standard errors, BLUPs, fit indices, and (planned) exported Hessians/standardization stay consistent across languages.

When to pick libsemx

- You need SEM + mixed effects in one model (latent factors plus random effects).
- You have non-Gaussian outcomes (binomial, Poisson/neg-bin, survival, ordinal) with random effects.
- You need flexible or custom covariance structures (multi-kernel, Kronecker, AR/Toeplitz).
- You want reproducible parity between Python and R for the same analysis.

When lme4/nlme might be faster

- Small Gaussian LMMs with simple random structures where `lme4`'s specialized code path dominates. For these, you can lower `libsemx` tolerances/iterations (`max_iterations = 50, tolerance = 1e-4`) to reduce overhead, but `lme4` will likely remain faster on tiny datasets. On larger or more complex models, the gap narrows because likelihood evaluation—not setup overhead—dominates runtime.

0.12 Interpreting fit indices and degrees of freedom

- Chi-square/df/CFI/TLI/RMSEA/SRMR are computed only when SEM sample statistics are available (latent-variable models or multi-group SEM). GLMM-only fits (e.g., survival, GBLUP, simple mixed models) leave `df` as `NA` and the chi-square carries no SEM meaning there. Use AIC/BIC/variance components/BLUPs for mixed models; use SEM indices only for latent-variable analyses with covariance structure.

- Bates, Douglas, Martin Mächler, Ben Bolker, and Steve Walker. 2015. “Fitting Linear Mixed-Effects Models Using Lme4.” *Journal of Statistical Software* 67 (1): 1–48. <https://doi.org/10.18637/jss.v067.i01>.
- Covarrubias-Pazaran, Giovanny. 2016. “Genome-Assisted Prediction of Quantitative Traits Using the r Package Sommer.” *PLOS ONE* 11 (6): e0156744. <https://doi.org/10.1371/journal.pone.0156744>.
- Seabold, Skipper, and Josef Perktold. 2010. “Statsmodels: Econometric and Statistical Modeling with Python.” In *9th Python in Science Conference*.